


# MacSim Tutorial Part-III

Jaekyu Lee



Georgia  
Tech



comparch



# Agenda (Part III)

- | Source code directory structure
- | Trace Generation
- | Memory System
- | Process Manager
- | Data structures
- | Knobs and Stats
- | Debugging



# Source Code

**Jaekyu Lee**



# MacSim Source Tree

MacSim Top-  
level

`bin/` Build output directory  
`def/` Contain definitions of parameter and events for  
statistics  
`doc/` Documentation (doxygen)  
`params/` Include sample parameter configuration files  
`scripts/` Scripts to build MacSim binary  
`src/` Source files (.cc and .h)  
`tools/` Tools (x86 trace generator, trace reader, ...)

// These files containing procedure to build and run MacSim and  
`SConstruct`, `Sconscript`, `buildy.py`

// macsim-sst  
`Makefile.am`  
`macsimComponent.cpp`  
`macsimComponent.h`



# Build Process Demo



# Doxygen Documentation

| Detailed MacSim class (members, hierarchy) information

| **Doxygen url:**

- <http://comparch.gatech.edu/hparch/macsim/doxygen/html/index.html>

| Details in:

- [macsim/doc/macsim.pdf](#)

Table 5. Source files and their purpose/content

File(s)	Purpose
frontend.cc/h, fetch_factory.cc/h, bp*.cc/h	Fetch stage
allocate*.cc/h, rob*.cc/h, map.cc/h	Decode and Allocate stages
schedule*.cc/h	Schedule stage
exec.cc/h	Execution stage
retire.cc/h	Retire stage
port.cc/h, cache.cc/h dram.cc/h, memory*.cc/h, memreq_info.cc/h, readonly_cache.cc/h, sw_managed_cache.cc/h	Memory system
pref*.cc/h	Prefetchers
trace_read.cc/h inst_info.h	Reading traces
core.cc/h	Class representing a core being simulated
process_manager.cc/h	Process Manager/thread scheduler
uop.cc/h	Uop structure and related enums
macsim.cc/h	Class containing pointers to the simulated cores, NoC,
memory system, knobs and other objects	
knob.cc/h	Classes for supporting knobs
statistics.cc/h	Classes for supporting knobs
factory_class.cc/h	Implementation of different factory classes
bug_detector.cc/h	Class useful for debugging forward progress errors happen
utils.cc/h	Utility classes and functions
debug_macros.h	Macros for debugging
assert_macros.h	Macros for assert statements with debug information
global*.h	Forward declarations and typedefs

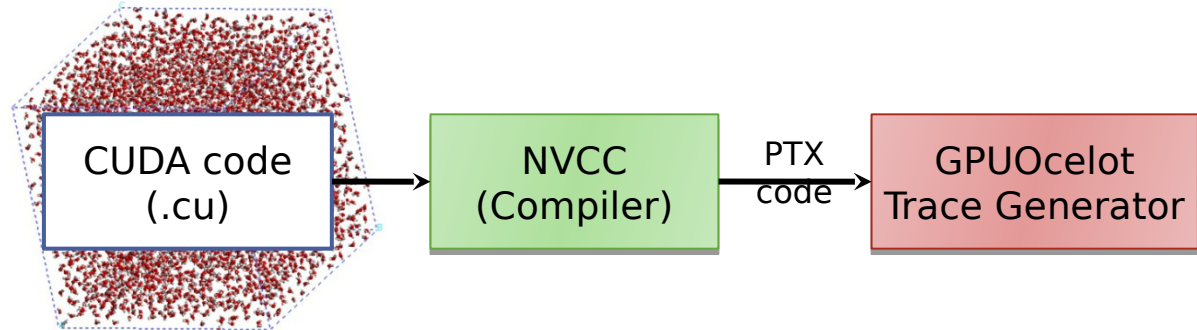


# Trace Generation




# Overview of Trace Generation


## | GPGPU traces



## | X86 traces



## | Graphics traces (ongoing)





# Trace Format



## | Trace has following fields

- PC, Opcode, Operands, Memory addresses, Memory request size, Active bits (to indicate divergent warp)

## | Each thread (warp) has own file with its thread id and

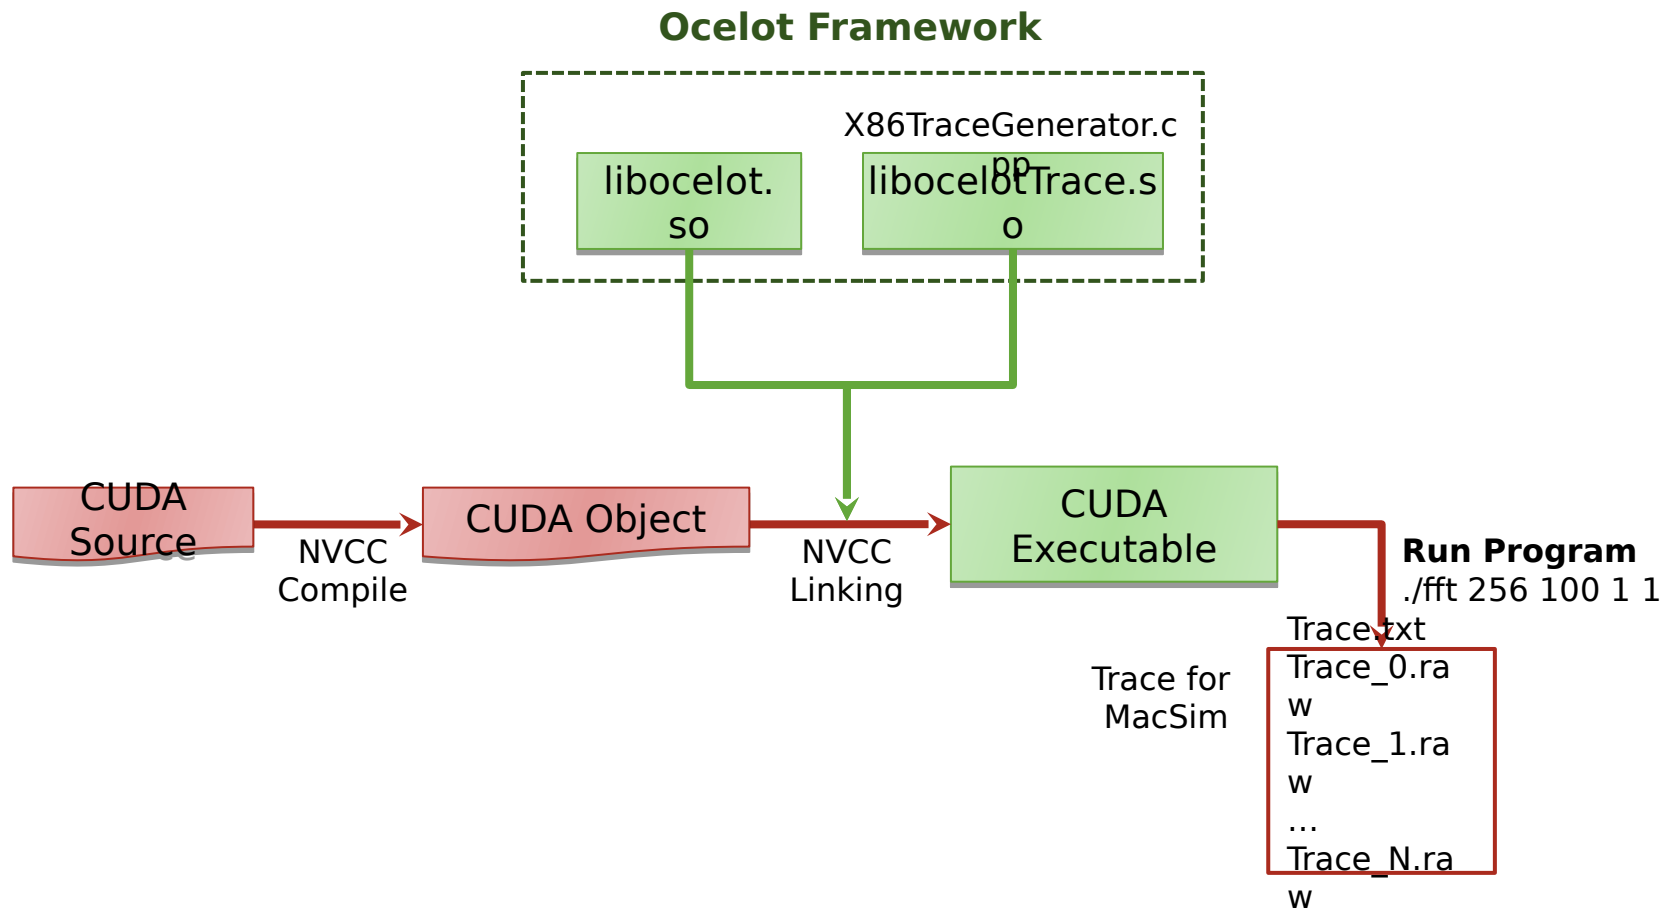
Trace_0.raw	TID: 0	BID: 0
Trace_1.raw	TID: 1	BID: 0
Trace_65536.raw	TID: 65536	BID: 1
Trace_65537.raw	TID: 65537	BID: 1

## | In MacSim, these trace streams are converted to internal RISC-style micro-ops (uop)



# Overview of GPU Trace Generation

MacSim uses **GPUOcelot** for GPU trace generation





# Steps to Generate GPU Traces

| Please refer to

- <http://code.google.com/p/gpuocelot/wiki/TraceGeneration> (Ocelot)
- <http://code.google.com/p/macsim/wiki/TraceGeneration> (MacSim)
- Documentation in MacSim repository

1. `svn checkout http://gpuocelot.googlecode.com/svn/trunk gpuocelot`  
Checkout GPUOcelot from subversion repository

2. `cd gpuocelot/ocelot; sudo ./build.py -install  
cd gpuocelot/trace-generators; libtoolize; aclocal; autoconf; automake;  
./configure; make; sudo make install`

# Steps to Generate GPU Traces -

## cnt'd

### 3. Link libocelot and libocelotTrace against CUDA applications

```
nvcc sourcefile.cu -locelot -ocelotTrace -lz -lboostxxx (boost libraries)
```

### 4.

# Steps to Generate GPU Traces -

## cnt'd

### | Setting environment variables

- TRACE\_PATH : directory to store traces (default: current dir)
- TRACE\_NAME : prefix name for traces (default: Trace)
- **KERNEL\_INFO\_PATH** : file that contains kernel information
- **COMPUTE\_VERSION** : compute capability

### | Example

In ~/.bashrc

```
# Create a trace directory in the /storage/traces
export TRACE_PATH="/storage/traces/"
```

```
# kernel_info has the kernel information
export KERNEL_INFO_PATH="kernel_info"
```

```
# Calculate occupancy based on compute capability 2.0
export COMPUTE_VERSION="2.0"
```

# Steps to Generate GPU Traces -

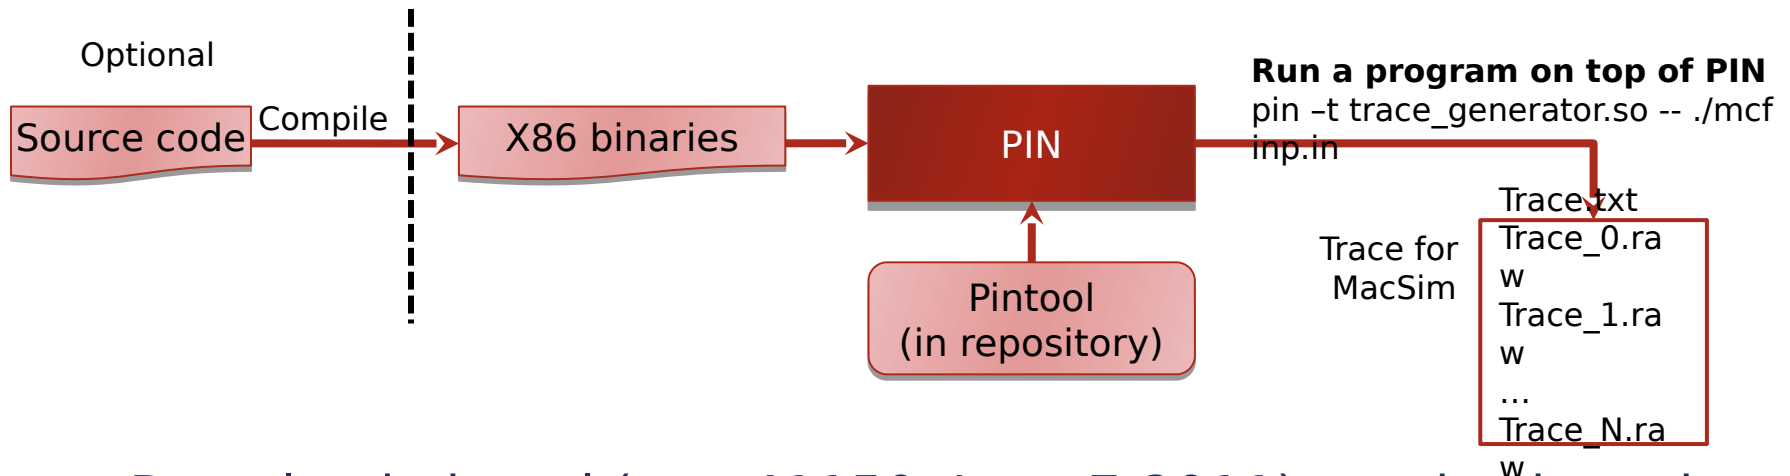
## cnt'd

- | The kernel information file (kernel\_info) should contain the following information.
  - kernel name
  - register usage (per thread)
  - shared memory usage (per thread)
  - (e.g.) `_Z9Memcpy_SWPfs_i 14 52`
- | Running CUDA application creates a directory with the kernel name where traces are generated as well as `kernel_config.txt` which is used for MacSim.

```
drwxr-xr-x  2 anonymous group 69632 Sep 00 22:03 _Z9Memcpy_SWPfs_i_0
-rw-r--r--  1 anonymous group   66 Sep 09 22:29 kernel_config.txt
```

# Steps to Generate CPU Traces

MacSim uses **Pin** for CPU trace generation



- Download pintool (ver. 41150, June-7-2011), not backward compatible
- Build pintool from source in repository  
`cd macsim/tools/trace_generator; make`
- Run `pin -t trace_generator.so [options] - binary input`

- **Pin homepage and tutorial:** <http://www.pintool.org>



# Micro-ops (uops)

- | Traces are translated to internal RISC-style micro-ops
  - All traces are shared same structure
  - Defined in `trace_read.h` and `trace_read.cc`
  
- | GPU traces: almost 1-1 matching
- | CPU traces: simple translation
  - information is coming from Pin's XED
  - XED provides `Inst_category`, register ids etc.
    - ▶ Does op have load/store?
    - ▶ Does op need a temp register?
    - ▶ Load+control flow, Load+store, load+no destination register
    - ▶ Does op have repeat?
    - ▶ Does op have a control flow instruction?





# MacSim Memory System

**Jaekyu Lee**

# Overview of Memory System of MacSim

## Cache hierarchy

- One or multiple levels
- Highly flexible

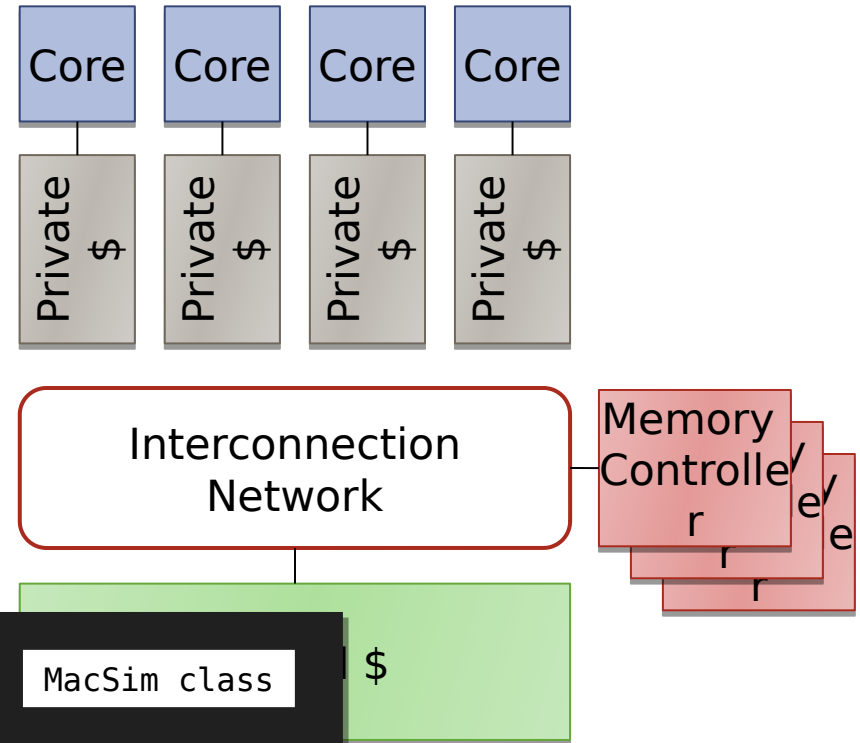
## Interconnection Network

- Default, IRIS
- Ring, 2D mesh and torus

## Memory Controller

- FCFS
- FR-FCFS

NoC and MCs are easily attachable/detachable with other modules



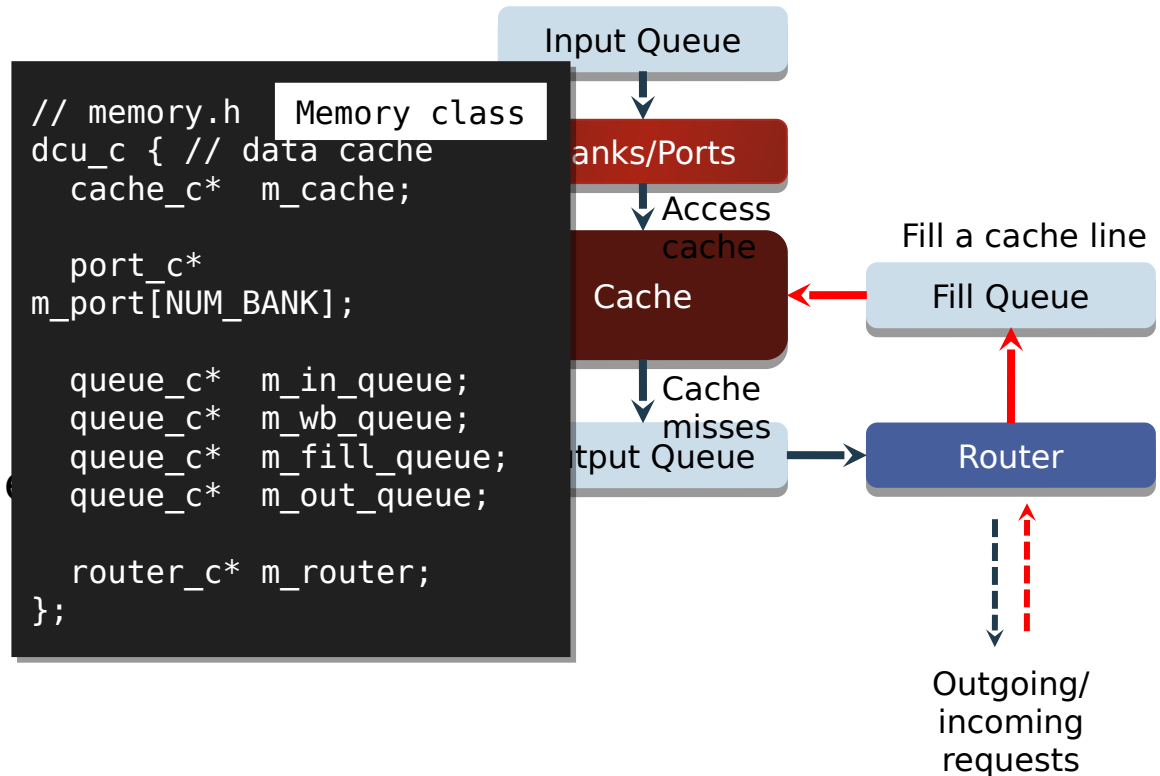
```
// macsim.h
macsim_c {
    memory_c* m_memory;
    dram_controller_c**
    m_dram_controller;
    router_wrapper_c* m_router;
};
```

MacSim class

# Cache

## | Each cache level consists

- Cache
  - ▶ Banked
  - ▶ Multiple ports
- Multiple queues
  - ▶ Input Queue
  - ▶ Output Queue
  - ▶ Write-back Queue
  - ▶ Fill Queue
- Router

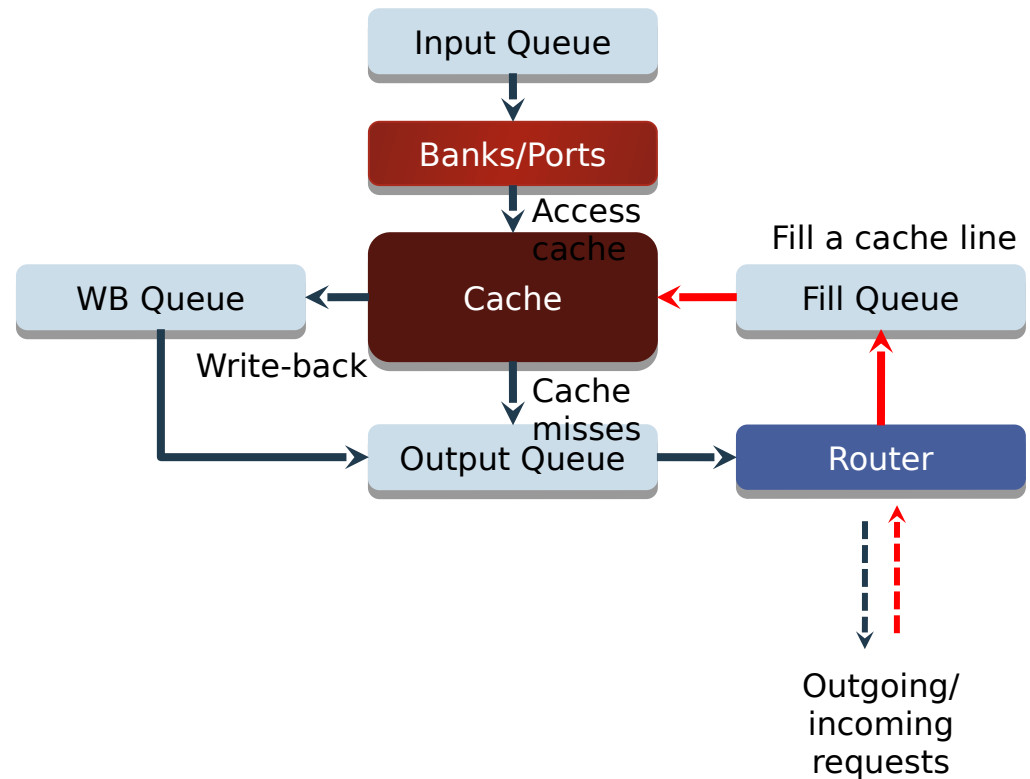




# Cache - cnt'd

## | Cache operations and flows between queues

- Cache access
- Cache eviction
- Outgoing requests
- Cache line fill






# How to Construct Cache Hierarchy

| To model various cache hierarchy models, we employ very flexible cache hierarchy

| We need to define

- Cache has a router
  - If router is enabled, traffics between upper/lower caches are thru the router
- Cache has a direct link between upper/lower cache
  - Otherwise, there must be a direct link between caches





# How to Construct Cache Hierarchy - cnt'd

```
// cache constructor
dcu_c(
    int id,                // cache id
    Unit_Type type,        // core type
    int level,             // level
    memory_c* mem,         // pointer to the memory system
    int noc_id,            // router id
    dcu_c** next,          // pointer to upward cache
    dcu_c** prev,          // pointer to downward cache
    macsim_c* simBase);    // pointer to simulation base class
```

Cache constructor

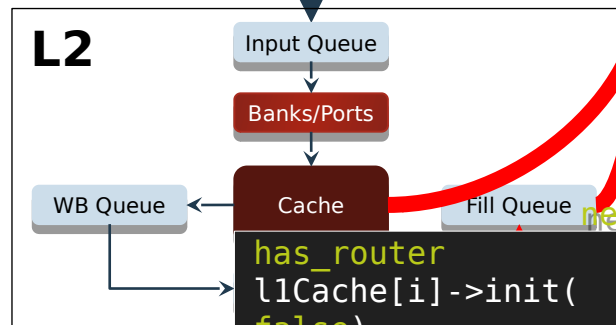

```
// cache initialization function
void dcu_c::init(
    int next_id,           // downward cache id
    int prev_id,           // upward cache id
    bool coupled_up,       // has direct link to upward cache
    bool coupled_down,     // has direct link to downward cache
    bool disable,          // cache disabled?
    bool has_router);      // has router?
```

Cache  
initialization



# Cache Hierarchy Example - 1

## 2-level private cache with the direct link



L1 does not have router

L1 miss

- Inserted directly into L2's input queue

L2 hit

- Insert data directly into L1's fill queue

L2 miss

- Go thru L2's router

next\_id, prev\_id, coupled up, coupled down, disable


```

has_router
l1Cache[i]->init( i, -1, false, true, false,
false)
l2cache[i]->init( -1, i, true, false, false,
true)
  
```



# Cache Hierarchy Example - 2

| Private L1, L2 caches with Shared L3 cache (No direct link)



| L2 miss


- Go thru L2's router

| L3 hit

- Go thru L3's router to feed data to L2

| L3 miss

- Go thru L3's router
- L2 fill - Go thru L3's router to feed data to L2



next\_id, prev\_id, coupled\_up, coupled\_down, disable,

```

has_router
l1cache[i]->init( i, -1, false, true, false,
false)
l2cache[i]->init( -1, i, true, false, false,
true)
l3cache[i]->init( -1, -1, false, false, false,
true)
  
```






# Cache Hierarchy Example - 3


## Private L1, L2 caches with Shared L3 cache (With direct link)

- Shared L3 has multiple tiles (bank)



### L2 miss

- Destination: local tile - Use direct link
- Destination: distant tile - Go thru L2's router



### L3 hit

- Destination: local tile - use direct link
- Destination: distance tile - Go thru L3's router to feed data to L2

### L3 miss

next\_id, prev\_id, coupled\_up, coupled\_down, disable,

has\_router

```
l1cache[i]->init( i, -1, false, true, false,
false)
l2cache[i]->init( -1, i, true, true, false,
true)
l3cache[i]->init( -1, i, true, false, false,
true)
```



# Cache Hierarchy - cnt'd

- | We provide sample cache hierarchies
  - Users need to specify single parameter to use sample classes
    - ▶ No cache - NVIDIA G80 GPU architecture
    - ▶ Private L1 / Shared L2 - NVIDIA Fermi GPU architecture, Intel Core architecture
    - ▶ Private L1, L2 / Shared+tiled L3 with direct link - Intel Sandy Bridge architecture
    - ▶ Private L1, L2 / Shared+tiled L3 without direct link - Generic 2D topology

Parameter: <string> KNOB\_MEMORY\_TYPE

Sample memory  
classes

```
class no_cache_c           // no cache architecture (NVIDIA G80)
class l3_coupled_network_c // 3-level, direct link (Intel Sandy Bridge)
class l3_decoupled_network_c // 3-level, no direct link (General 2D
topology)
class l2_decoupled_network_c // 2-level, no direct link (NVIDIA Fermi)
```



# Interconnection Network

- | Default bi-direction ring network
  - Pipelined router architecture
  - Model virtual and physical channels with arbitrations
  
- | IRIS interconnection simulator
  - CASL research group (Georgia Tech, led by Prof. Yalamanchili)
  - Model more detailed features
  - Have various topologies (Ring, 2D Mesh, Torus, ...)



# Memory Controller

- | We model DRAM banks and channels
- | DRAM bandwidth is modeled
- | 3 timing parameters
  - Activate, Precharge, Column Access
- | Scheduling policies
  - FCFS, FR-FCFS
  - Other policies can be easily implementable

```
// dram.h
class dram_controller_c {
    address_parsing
    dram_request_buffer,
    banks,
    channels,
    router,
    ...
};
```

DRAM controller  
pseudo code



# Memory Access Handling

## | Request Merging

- Memory requests with the same address can be merged with other older requests in various locations
- Intra-core: MSHR
- Inter-core: DRAM request buffer

## | Coalesced vs. Uncoalesced

- If GPU SIMD threads generate memory requests contiguous address, multiple requests can be coalesced in a larger request
- We model coalescing in Trace Generator
- GPUOcelot provides memory address and access size for each thread of a warp (block)



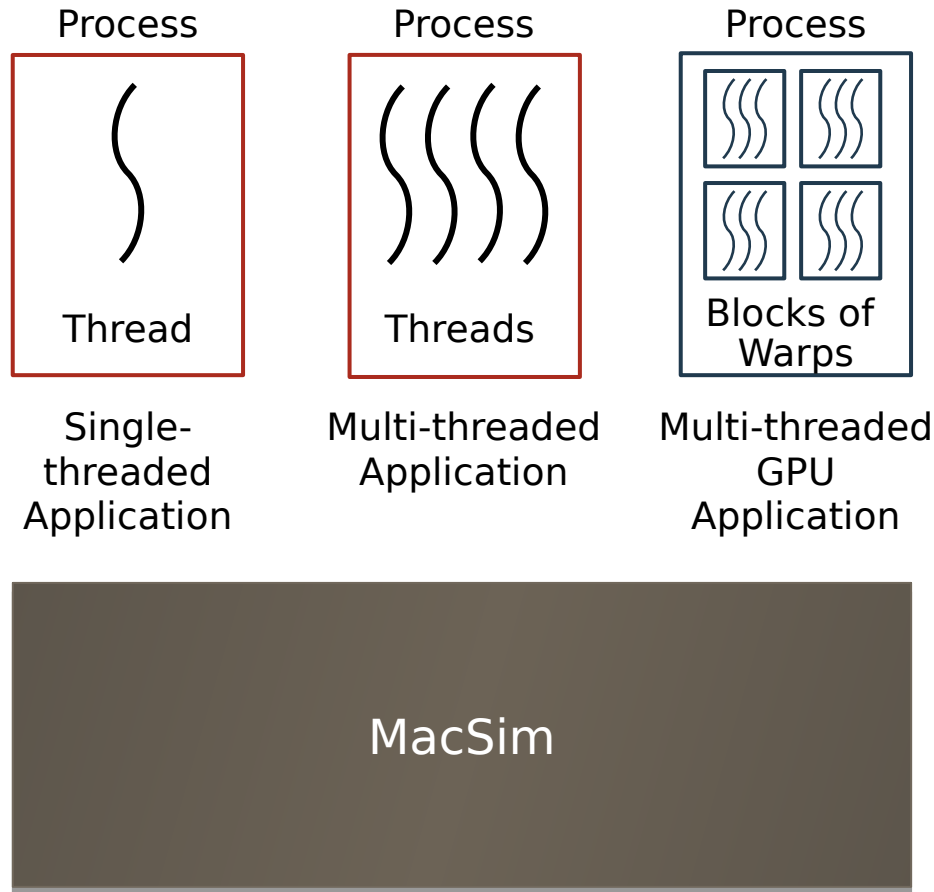
# Process Manager



# Process Manager

- | We use term the “Process” for an application in MacSim
- | A Process consists of one or more threads (thread blocks)
- | Process Manager is the component of MacSim that assigns work to cores
- | It also handles creation/termination of process/threads (thread blocks)
  - Bookkeeping to track number of running threads, completed threads

# Process Manager Capabilities



## Single process simulation

- Single-threaded CPU
- Multi-threaded CPU
- GPGPU

## Multi-process simulation

- Any combination of processes

## Configuring Multi-process simulations

- Partition cores among applications
- Specify the number of threads (thread blocks) per core
- Specify repetition flag





# Process/Thread Creation

- | Allocate process structure, read trace config file and determine when threads have to be created
  - main thread created at start up
  - points of creation of child threads specified in terms of number of instructions executed by main thread
    - as the main thread executes check if a child thread should be created, if so, create the child thread and mark it ready for execution
  - for GPGPU applications all warps are created and marked ready at start up



# Process/Thread Creation

On application start up

`create_process()`

`read_trace()`  
in `trace_read.cc`

During execution main  
thread spawns child threads

`create_thread_node()`

Create dummy thread node  
for scheduling. Actual  
creation of thread (memory  
allocation, ..) happens when  
a thread is scheduled to a  
core

`insert_thread()/`  
`insert_block()`

Insert dummy node to pending  
queue  
data  
structures

`m_thread_queue /`  
`m_block_queue`



# Thread (Thread Block) Scheduling

## | CPU

- Greedy assignment of threads to cores
- Each core assigned threads up to user specified limit

## | GPU

- Scheduling at thread block granularity - all warps in a block assigned to same core
- Greedy assignment of blocks to cores (other policies can be added easily)
- Occupancy - How many blocks per core?
  - Calculated using Ocelot at trace generation and passed to Mascim via traces
  - Can be overridden by user

| Once assigned threads/thread block remain on the same core until termination

| Multi-application simulations (CPU+CPU or CPU+GPU),


- Each application is restricted to a statically defined set of cores
- Can repeat short running applications until a long running application ends




# Thread Scheduling

- | Process Manager maintains queues of ready but unassigned threads/thread blocks

X86 Queue



GPU Queue



max threads per core  
= 2

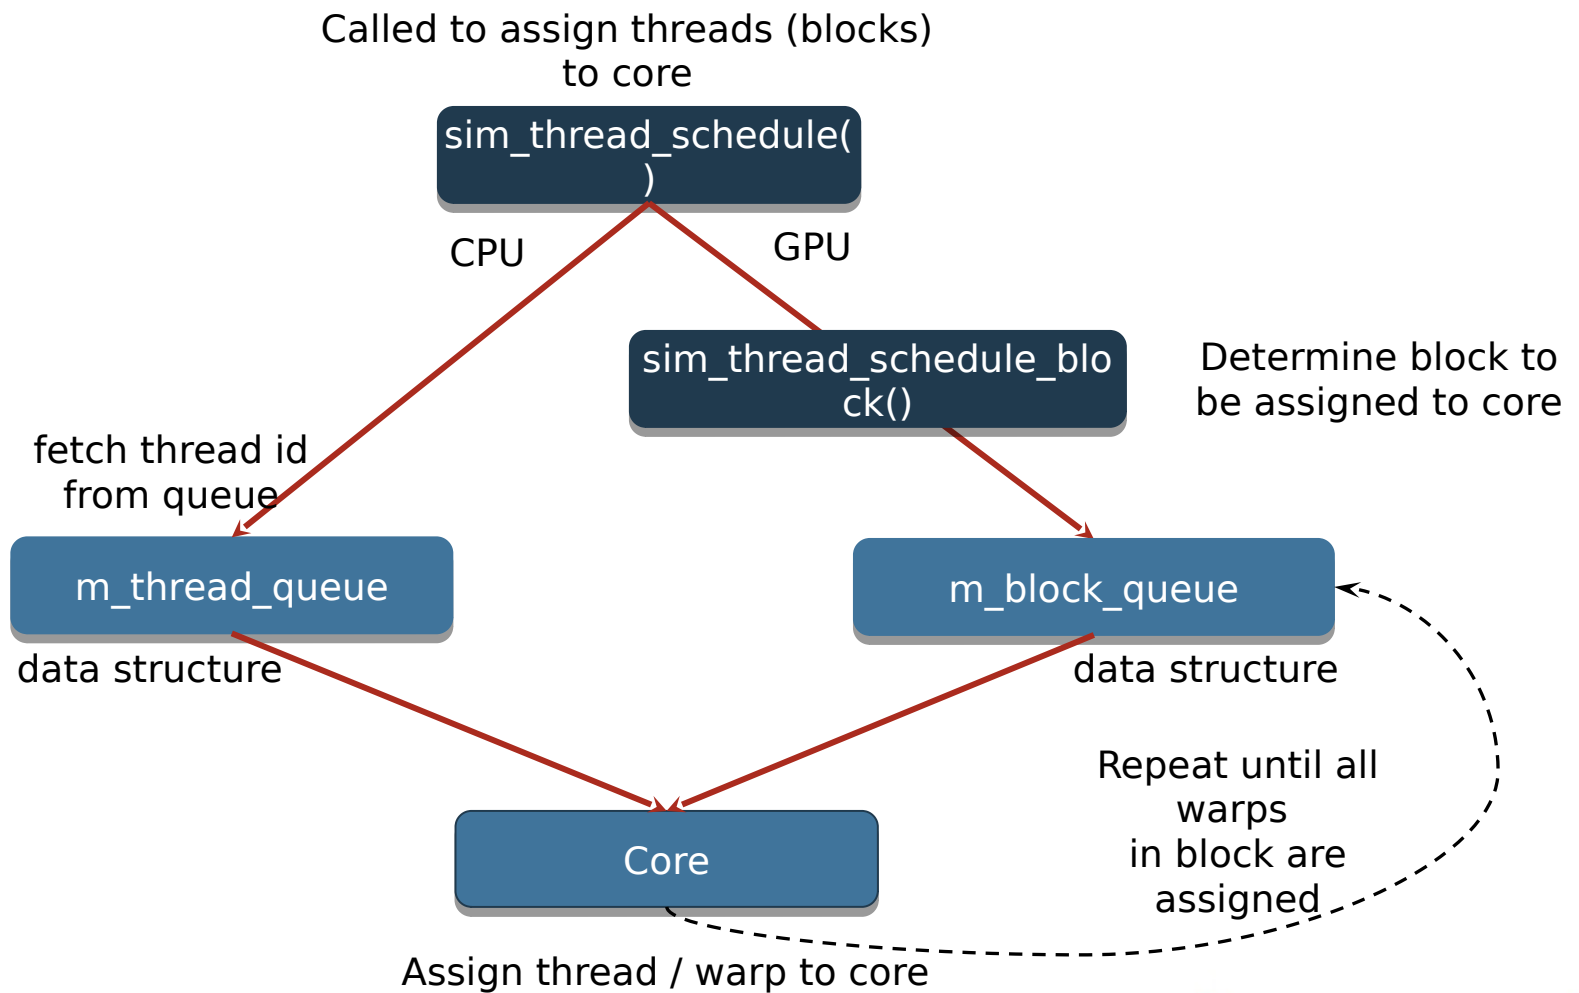
- | Initially - greedy assignment to cores until specified limit
- | Thereafter - assign a new thread to a core whenever a thread on the core terminates



max thread blocks per core  
= 2

- | Initially - greedy assignment to cores until specified limit
- | Thereafter - assign a new thread block to a core only when the core retires an entire thread block

# Thread Scheduling





# Process Manager

- | Consists of two files
  - `process_manager.h/cc`
  
- | `process_manager.h`
  - Declares Process Manager class
    - `process_manager_c`
  - Defines data structures for
    - Applications being simulated
      - `process_s`
    - Warps on GPU cores / threads on CPU cores
      - `thread_s`
    - Bookkeeping
      - `block_schedule_info_s`
      - `thread_trace_info_node_s`



# Process Manager

| `process_manager.cc`

- Defines member functions of `process_manager_c` class



# Process Manager

## | thread\_s

- structure analogous to task structures used in OS kernels (task\_struct in Linux)
- each CPU thread and GPU warp has an instance of thread\_s
- contains fields for
  - ▶ thread id, block id, pointer to trace file, process to which thread/warp belongs, instructions read, trace buffer
  - ▶ In GPU cores, thread id represents warp id





# Process Manager

## | thread\_s

```
typedef struct thread_s {
    int          m_thread_id; /**< current thread id */
    int          m_block_id; /**< block id */
    gzFile       m_trace_file; /**< gzip trace file */
    bool         m_file_opened; /**< trace file opened? */
    bool         m_main_thread; /**< main thread (usually thread id 0)
*/
    uint64_t     m_inst_count; /**< total instruction counts */
    uint64_t     m_uop_count;  /**< total uop counts */
    bool         m_trace_ended; /**< trace ended */
    process_s*   m_process; /**< point to the application belongs to */
    bool         m_ptx; /**< GPU thread */
    char*        m_buffer; /**< trace buffer */
    frontend_s*  m_fetch_data; /**< frontend fetch data */
    trace_info_s* m_prev_trace_info; /**< prev instruction trace info */
    trace_info_s* m_next_trace_info; /**< next instruction trace info */
    bool         m_thread_init; /**< thread initialized */
    trace_uop_s* m_trace_uop_array[MAX_PUP]; /**< trace uop array */
    trace_uop_s* m_next_trace_uop_array[MAX_PUP]; /**< next trace uop
array */
} thread_s;
```



# Process Manager

## | process\_s

- one instance for each application being simulated
- common structure for both CPU and GPGPU applications
  - ▶ GPGPU applications do not include CPU-side traces, they only include traces from all kernels executed by the application
- has fields for process id, start info of threads, no. of threads (warps) in application, threads (warps) created, threads terminated (warps), list of valid cores for this application, list of kernels executed by application (GPGPU only)

# Process Manager



## process\_s

```
typedef struct process_s {
    process_s();
    ~process_s();

    unsigned int    m_process_id; /**< current process id */
    int             m_orig_pid; /**< original process id - in case of repetition */
    int             m_max_block; /**< max blocks per core for the application */
    thread_start_info_s *m_thread_start_info; /**< thread start information */
    thread_s**      m_thread_trace_info; /**< thread trace information */
    unsigned int    m_no_of_threads; /**< number of total threads */
    unsigned int    m_no_of_threads_created; /**< number of threads created */
    unsigned int    m_no_of_threads_terminated; /**< number of terminated threads */
    map<int, bool>   m_core_list; /**< list of cores that this process is executed */
    queue<int>      *m_core_pool; /**< core pool pointer */
    bool            m_ptx; /**< GPU application */
    int             m_repeat; /**< application has been re-executed */
    vector<string>   m_applications; /**< list of sub-applications */
    vector<int>      m_kernel_block_start_count; /**< block id start count for sub-
appl. */
    string          m_current_file_name_base; /**< current sub-appl.'s filename base
*/
    string          m_kernel_config_name; /**< kernel config file name */
    int             m_current_vector_index; /**< current index to the sub-
application */
    map<int, bool>   m_block_list; /**< list of block currently running */
    uns64           m_inst_count_tot; /**< total instruction counts */
    int             m_block_count; /**< total block counts */
} process_s;
```



# Process Manager

## | block\_schedule\_info\_s

- bookkeeping structure for thread blocks in a GPGPU application, one instance for each thread block
- contains fields for no. of warps in block, no. of terminated warps, core to which block is assigned

```
typedef struct block_schedule_info_s {  
    bool    m_start_to_fetch;    /**< start fetching */  
    int     m_dispatched_core_id; /**< core id in which this block is launched  
*/  
    bool    m_retired;           /**< retired */  
    int     m_dispatched_thread_num; /**< number of dispatched threads */  
    int     m_retired_thread_num; /**< number of retired threads */  
    int     m_total_thread_num;  /**< number of total threads */  
    int     m_dispatch_done;     /**< dispatch done */  
    bool    m_trace_exist;       /**< trace exist */  
    Counter m_sched_cycle;       /**< scheduled cycle */  
    Counter m_retire_cycle;      /**< retired cycle */  
} block_schedule_info_s;
```



# Process Manager

## | thread\_trace\_info\_node\_s

- wrapper structure around thread\_s used by process\_manager\_c to track threads (warps) yet to be assigned to cores

```
typedef struct thread_trace_info_node_s {  
    thread_s* m_trace_info_ptr; /**< trace information pointer */  
    process_s* m_process;      /**< pointer to the process */  
    int m_tid;                 /**< thread id */  
    bool m_main;               /**< main thread */  
    bool m_ptx;                /**< GPU simulation */  
    int m_block_id;            /**< block id */  
} thread_trace_info_node_s;
```



# Process Manager

## | process\_manager\_c

- creates processes, threads (warps)
- does assignment of threads / thread blocks to cores
- invoked again on termination of threads (warps)
- contains list of threads not assigned to any core (CPU only)
- contains list of threads blocks not assigned to any core (GPGPU only)
- handles repetition of applications (triggered by retire) when multiple applications are being run



# Process Manager

## Process/Thread creation and termination

```
// allocates process's node, reads config file and determines  
// no. of threads/warps and thread blocks in application / each  
kernel  
// of application, start info of threads  
create_process()  
  
// called for each thread/warp when it becomes ready for  
launch,  
// allocates a thread_trace_info_node_s node and inserts into  
// m_thread_queue or m_block_queue. All warps are ready for  
launch  
// at the start of a kernel  
create_thread_node()  
  
// allocates and initializes a thread_s node, opens trace  
file  
// for thread/warp  
create_thread()  
  
// cleans up some data structures and saves stats (if this  
the  
// first run of the application)  
terminate_process()  
  
// clears data structures, retires thread block if all warps  
in  
// block have completed, calls sim_thread_schedule()  
terminate_thread()
```



# Process Manager

## Thread/block scheduling

```
// assigns a thread to a core or a thread block to a core if
the
// number of threads/blocks on the core are fewer than the
maximum
// allowed
void process_manager_c::sim_thread_schedule(void);

// called by sim_thread_schedule() to determine the next
block
// that will be assigned to a core
int process_manager_c::sim_thread_schedule_block(int
core_id);

// inserts thread / thread_block into list of threads /
blocks yet
// to be launched
void
process_manager_c::insert_thread(thread_trace_info_node_s*)
void
process_manager_c::insert_block(thread_trace_info_node_s*)
```





# Classes/Data Structures

**Nagesh Lakshminarayana**



# Data Structures

## | uop\_c

- main data structure for uops
  - an instance is allocated for each uop
- contains thread (warp) id, core id, block id (GPU only), timestamps of when uop passed through different stages, opcode type, source and destination operands and all other information required for execution
  - contains list of child uops in case of uncoalesced memory accesses

| New uops are allocated from a uop pool and completed uops are returned to the pool



# Data Structures

## | core\_c

- class for cores in simulation
- contains pointers to objects for pipeline stages and components
- during simulation `run_a_cycle()` of `core_c` is called every cycle
  - ▶ `core_c::run_a_cycle()` calls `run_a_cycle()` function of the pipeline stages



# Data Structures

## | map\_c

- Class for tracking data and memory dependences between uops of a thread
- once a dependence between uops is identified, pointer to source uop is added to the list of uops on which the dependent will wait
  - in the schedule stage a check is made to see if all sources of a uop have completed



# Data Structures

## | mem\_req\_s

- data structure for memory requests
  - ▶ an instance of mem\_req\_s is allocated for every memory request
- contains id information, details of primary memory request and piggybacking requests, pointer to done function
  - ▶ done function is a callback function called when the memory request is completed (actually when response reaches L2)
    - user can add code here for collecting statistics



# Data Structures

## | drb\_entry\_s

- data structure for requests in DRAM
  - ▶ contains timestamp, row id, bank id and pointer to memory request
- requests to a bank are held in a list, DRAM scheduler sorts this list according to policy being implemented and returns the head of the list
  - ▶ can extend the structure to track any information that your DRAM policy requires



# Utility Classes

**Nagesh Lakshminarayana**



# Utility Classes

## | pool\_c

- class used to implement memory pools of objects/structures of different types
  - `acquire_entry()` and `release_entry()` for getting/releasing an from/to the pool
  - pool expands automatically when it is out of space

## | pqueue\_c

- class for modeling latencies of pipeline stages
  - entries pushed into the queue are ready to be removed from queue after the specified latency
  - priority used among entries pushed into queue in the same cycle





# Utility Classes

## | Factory classes

- Fetch policies, DRAM policies, memory hierarchies and so on are implemented by classes whose objects are instantiated via factory mechanisms
  - ▶ One factory class for fetch policies, one for DRAM policies, ...
- For supporting new policies
  - ▶ Define a class for the policy
  - ▶ Register class with factory under the name of the policy
  - ▶ Specify new policy as the policy to be used during simulation (via simulation parameters)



# Knob Variables

**Nagesh Lakshminarayana**



# Knobs

- | Used to specify architecture and simulation parameter values
- | Available knobs can be found in .param.def files in def/ directory under main source directory
- | To set a knob
  - provide a value on the command line (highest priority),
  - provide a value in the parameter file, or
  - provide a default value in the definition (lowest priority)
- | All knobs and their values for a simulation are written to *params.out* file when simulation starts



# Adding New Knobs

- | Add a definition in a relevant .param.def file in def/ directory (create a new file if necessary)
  - Knob definitions converted into c++ source and included in build process automatically
  
- | Format of definition
  - `param<{name_in_code}, {name}, {knob_type}, {default_value}>`
    - `name_in_code` – name for accessing knob in the code (in upper case)
    - `name` – name used in command line or parameter file (in lower case)
    - `knob_type` – bool, different int types, float, string
    - default value must be provided
  
- | Example definition
  - `param<L2_ASSOC, l2_assoc, int, 8>`



# Setting and Using Knobs

## | Example definition

- `param<L2_ASSOC, l2_assoc, int, 8>`

## | Setting value on command line

- `./macsim --l2_assoc=16`

## | Setting value in parameter file (params.in)

- `l2_assoc 16`

## | Accessing in code

- use `KNOB_L2_ASSOC` (prefix knob name with `KNOB_`)
  - `*m_simBase->m_knobs->KNOB_L2_ASSOC`, or
  - `m_simBase->m_knobs->KNOB_L2_ASSOC->getValue()`



# Statistics

**Nagesh Lakshminarayana**



# Stats

- | Global or per core stats
- | Three Stat types supported
  - COUNT
    - #occurrences of an event
  - RATIO
    - Ratio of #occurrences of event A to the #occurrences of event B
  - DIST
    - Proportion of each event in a group of events
- | Stat definitions can be found in .stat.def files in def/ directory under main source directory



# Stats

- | At the end of simulation .stat.out files containing stat values are generated
  - single application simulation - .stat.out files
  - Multi-application simulation - .stat.out.<appl\_id> files
    - stat files generated when the first run of application completes
  
- | Script to read stat values from the stat files will be provided
  - provide name of stat and value of stat will be printed
  - script can also evaluate simple expressions containing stat names





# COUNT Stat

## Format of definition

- `DEF_STAT(STAT_NAME, COUNT, NO_RATIO [, PER_CORE])`
  - `STAT_NAME` - name for accessing stat in code (in upper case)
  - `NO_RATIO` - not a ratio stat
  - `PER_CORE` - optional, use if you want a per core stat

## Example definitions

- `DEF_STAT(INST_COUNT_TOT, COUNT, NO_RATIO)`
- `DEF_STAT(INST_COUNT, COUNT, NO_RATIO, PER_CORE)`

## Example output

- |                                |        |        |
|--------------------------------|--------|--------|
| <code>INST_COUNT_TOT</code>    | 100000 | 100000 |
| <code>INST_COUNT_CORE_0</code> | 52342  | 52342  |
| <code>INST_COUNT_CORE_1</code> | 47658  | 47658  |

Count



# RATIO Stat

## | Format of definition

- `DEF_STAT(STAT_NAME, RATIO, BASE_STAT_NAME [, PER_CORE])`
  - `STAT_NAME` - name for accessing stat in code (in upper case)
  - `BASE_STAT_NAME` - stat whose value is the denominator in the ratio
  - `PER_CORE` - optional, use if you want a per core stat
  - Requires the definition of a stat of type COUNT

## | Example definition

- `DEF_STAT(DCACHE_ACCESSES, COUNT, NO_RATIO)`  
`DEF_STAT(DCACHE_HITS, RATIO, DCACHE_ACCESSES)`

## | Example output

- `DCACHE_ACCESSES`  
`DCACHE_HITS`

67258	67258
8434	0.123
Count	

Count

Ratio = (8434 /  
 67258) =  
 (#DCACHE\_HITS/#DCACHE\_ACCESSES)

# DIST Stat



## | Format of definition

- `DEF_STAT(STAT_NAME_START, DIST, NO_RATIO [, PER_CORE])`  
`DEF_STAT(STAT_NAME, COUNT, NO_RATIO [, PER_CORE])*`  
`DEF_STAT(STAT_NAME_END, DIST, NO_RATIO [, PER_CORE])`
  - DIST stat should contain at least two stats

## | Example definition

- `DEF_STAT(ICACHE_HITS, DIST, NO_RATIO)`  
`DEF_STAT(ICACHE_MISSES, DIST, NO_RATIO)`

## | Sample output

- |               |        |        |
|---------------|--------|--------|
| ICACHE_HITS   | 292331 | 0.8937 |
| ICACHE_MISSES | 34767  | 0.1063 |

Proportion =  $(292331 / (292331 + 34767)) = 89.37\%$

Proportion =  $(34767 / (292331 + 34767)) = 10.63\%$

Count

$\Sigma(\text{prop}_i) = 1 =$   
100%



# Updating Stats

## | Macros for updating Stats

- Global stats
  - ▶ `STAT_EVENT(STAT_NAME)` - increment `STAT_NAME` by 1
  - ▶ `STAT_EVENT_N(STAT_NAME, n)` - increment `STAT_NAME` by `n`
  - ▶ `STAT_EVENT_M(STAT_NAME)` - decrement `STAT_NAME` by 1
- Per-core stats
  - ▶ `STAT_CORE_EVENT(CORE_ID, STAT_NAME)` - increment `STAT_NAME` for core `CORE_ID` by 1
  - ▶ `STAT_CORE_EVENT_N(CORE_ID, STAT_NAME, n)` - increment `STAT_NAME` for core `CORE_ID` by `n`
  - ▶ `STAT_CORE_EVENT_M(CORE_ID, STAT_NAME,)` - decrement `STAT_NAME` for core `CORE_ID` by 1



# Debugging

**Nagesh Lakshminarayana**



# Debug Statements

- | Debug statements can be printed only with the debug version of the MacSim binary
  - Run “make dbg” to build the debug version
  
- | To output debug statements for a stage/component
  - Set knob `debug_cycle_start` to the cycle value from when debug statements should be printed
    - Cycle value should be greater than zero
  - Set the debug enable knob for the stage/component to 1
    - Knobs available for pipeline stages, memory system and so on
    - Check `debug.param.def` for list of available knobs
  
- | Knob `debug_cycle_stop` can be set to turn off debug statements after some point
  - Set `debug_cycle_stop` to zero (default) to print till end of simulation



# Debug Statements

| params.in file with debugging enabled for front-end stage

- ...  
debug\_cycle\_start 1  
debug\_cycle\_stop 0  
debug\_front\_stage 1  
...



# Adding New Debug Statements

- | Debug statements can be added using the DEBUG macro
- | DEBUG macro takes parameters similar to the printf() function
- | Example debug statement
  - `DEBUG("m_core_id:%d thread_id:%d uop_num:%lld operands are not ready \n",  
m_core_id, cur_uop->m_thread_id, cur_uop->m_uop_num);`





# Sample Debug output

```
../src/frontend.cc:836: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 m_running_thread_num:1 m_fetching_thread_num:1 m_unique_scheduled_thread_num:1
../src/frontend.cc:914: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 try_again:1 tid:0
../src/frontend.cc:230: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 frontend fetch thread is:0
../src/frontend.cc:633: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 fetch_addr:46c856 new fetch_addr:46c856 m_icache hit
../src/frontend.cc:465: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): cycle_count:1136 m_core_id:0 tid:0 uop_num:60 inst_num:50 uop.va:
b54658 iaq:0 mem_type:2 dest:0 num_dests:0
../src/frontend.cc:550: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 tid:0 MT_scheduler[0]->0x46c85d
../src/frontend.cc:721: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 tid:0 inst_num:50 uop_num:60 opcode:11 isitEOM:1 sent
to qfe
../src/frontend.cc:465: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): cycle_count:1136 m_core_id:0 tid:0 uop_num:61 inst_num:51 uop.va:
0 iaq:0 mem_type:0 dest:19 num_dests:1
../src/frontend.cc:550: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 tid:0 MT_scheduler[0]->0x46c867
../src/frontend.cc:721: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1136): m_core_id:0 tid:0 inst_num:51 uop_num:61 opcode:11 isitEOM:1 sent
to qfe
../src/frontend.cc:836: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 m_running_thread_num:1 m_fetching_thread_num:1 m_unique_scheduled_thread_num:1
../src/frontend.cc:914: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 try_again:1 tid:0
../src/frontend.cc:230: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 frontend fetch thread is:0
../src/frontend.cc:633: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 fetch_addr:46c867 new fetch_addr:46c867 m_icache hit
../src/frontend.cc:465: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): cycle_count:1137 m_core_id:0 tid:0 uop_num:62 inst_num:52 uop.va:
64203e8 iaq:0 mem_type:2 dest:0 num_dests:0
../src/frontend.cc:550: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 tid:0 MT_scheduler[0]->0x46c86b
../src/frontend.cc:721: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 tid:0 inst_num:52 uop_num:62 opcode:11 isitEOM:1 sent
to qfe
../src/frontend.cc:465: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): cycle_count:1137 m_core_id:0 tid:0 uop_num:63 inst_num:53 uop.va:
64203f0 iaq:0 mem_type:2 dest:0 num_dests:0
../src/frontend.cc:550: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 tid:0 MT_scheduler[0]->0x46c86f
../src/frontend.cc:721: *m_simBase->m_knobs->KNOB_DEBUG_FRONT_STAGE (I=48 C=1137): m_core_id:0 tid:0 inst_num:53 uop_num:63 opcode:11 isitEOM:1 sent
to qfe
```



# Forward Progress Error

- | A Forward Progress Error is triggered when an active core has not retired any instructions for the number of cycles specified by the knob `forward_progress_limit`
  - A Forward Progress Error causes simulation to be aborted
  
- | A Forward Progress Error could be triggered when
  - A uop cannot be retired because the memory request it generated was lost in the memory system
  - A uop cannot be scheduled because one of its sources operands is not marked ready yet, but the uop producing the source has already retired (or is lost)



# Bug Detector

- | The knob `bug_detector_enable` can be turned on to generate additional debug information that would be helpful in resolving forward progress errors
  - Works with both debug and optimized versions of binary
  - Generates
    - ▶ `bug_detect_uop.out`
      - Dump of all uops in the pipeline of the failing core
    - ▶ `bug_detect_mem.out`
      - Dump of all memory requests issued by the failing core
    - ▶ `bug_detect_dram.out`
      - Dump of all requests in DRAM



# Questions?